

W271-2 – Spring 2016 – HW 2

Juanjo Carin, Kevin Davis, Ashley Levato, Minghu Song

February 10, 2016

Contents

Data	1
Exercises	2
Question 1	2
Question 2	6
Question 3	8
Question 4	10
Question 5	11
Question 6	12
Question 7	14
Question 8	14

Data

In the United States, a 401K is a type of retirement savings plan that is tied to a worker's place of employment. Employees that put money into a 401K enjoy certain tax benefits. Moreover, many employers have a policy of promoting 401K use, by matching some percentage of an employee's contributions. If an employer matches at, say, 50%, for every dollar that an employee puts into a 401K, the employer will put in another 50 cents.

The file `401k_w271.RData` contains data on 401k contributions that were filed with the IRS on form 5500. It was collected by Professor L. E. Papke and may have been further modified by the instructors to test your proficiency.

Exercises

Complete the following exercises, following the best practices outlined in class. Place your answers in a written report (pdf, word, or jupyter notebook format) along with relevant R statements and output.

Load the `401k_w271.RData` dataset and look at the value of the function `desc()` to see what variables are included.

```
load("401k_w271.Rdata")
```

Question 1

Your dependent variable will be `prate`, representing the fraction of a company's employees participating in its 401k plan. Because this variable is bounded between 0 and 1, a linear model without any transformations may not be the most ideal way to analyze the data, but we can still learn a lot from it. Examine the `prate` variable and comment on the shape of its distribution.

```
# Descriptive statistics of the whole dataset
desc
```

```
##   variable                                label
## 1   prate      participation rate, percent
## 2   mrate      401k plan match rate
## 3   totpart     total 401k participants
## 4   totelg     total eligible for 401k plan
## 5   age        age of 401k plan
## 6   totemp    total number of firm employees
## 7   sole = 1 if 401k is firm's sole plan
## 8   ltotemp    log of totemp
```

```
str(data)
```

```
## 'data.frame':   1534 obs. of  8 variables:
##  $ prate   : num  26.1 100 97.6 100 82.5 ...
##  $ mrate   : num  0.21 1.42 0.91 0.42 0.53 ...
##  $ totpart : num  1653 262 166 257 591 ...
##  $ totelg  : num  6322 262 170 257 716 ...
##  $ age     : int   8 6 10 7 28 7 31 13 21 10 ...
##  $ totemp  : num  8709 315 275 500 933 ...
##  $ sole    : int   0 1 1 0 1 1 1 0 1 1 ...
##  $ ltotemp : num   9.07 5.75 5.62 6.21 6.84 ...
## - attr(*, "datalabel")= chr ""
## - attr(*, "time.stamp")= chr "25 Jun 2011 23:03"
## - attr(*, "formats")= chr  "%7.0g" "%7.0g" "%7.0g" "%7.0g" ...
## - attr(*, "types")= int   254 254 254 254 251 254 251 254
## - attr(*, "val.labels")= chr  "" "" "" "" ...
## - attr(*, "var.labels")= chr  "participation rate, percent" "401k plan match rate" "total 401k part.
## - attr(*, "version")= int  10
```

```
summary(data)
```

```
##      prate      mrate      totpart      totelg
## Min.   : 3.00   Min.   :0.0100   Min.   : 50.0   Min.   : 51.0
## 1st Qu.: 78.10   1st Qu.:0.3000   1st Qu.: 156.2   1st Qu.: 176.0
## Median : 95.70   Median :0.4600   Median : 276.0   Median : 330.0
## Mean   : 87.56   Mean   :0.7315   Mean   : 1354.2   Mean   : 1628.5
## 3rd Qu.:100.00   3rd Qu.:0.8300   3rd Qu.: 749.5   3rd Qu.: 890.5
## Max.   :200.00   Max.   :4.9100   Max.   :58811.0   Max.   :70429.0
##      age      totemp      sole      ltotemp
## Min.   : 4.00   Min.   : 58   Min.   :0.0000   Min.   : 4.060
## 1st Qu.: 7.00   1st Qu.: 261   1st Qu.:0.0000   1st Qu.: 5.565
## Median : 9.00   Median : 588   Median :0.0000   Median : 6.377
## Mean   :13.18   Mean   : 3568   Mean   :0.4876   Mean   : 6.686
## 3rd Qu.:18.00   3rd Qu.: 1804   3rd Qu.:1.0000   3rd Qu.: 7.498
## Max.   :51.00   Max.   :144387   Max.   :1.0000   Max.   :11.880
```

```
# Descriptive statistics of prate
summary(data$prate)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##      3.00  78.10   95.70   87.56 100.00   200.00
```

```
round(stat.desc(data$prate, desc = TRUE, basic = TRUE, norm = TRUE), 2)
```

```
##      nbr.val  nbr.null  nbr.na      min      max
##      1534.00      0.00      0.00      3.00     200.00
##      range      sum      median      mean  SE.mean
##      197.00 134314.70      95.70      87.56      0.44
## CI.mean.0.95      var      std.dev      coef.var      skewness
##      0.87      300.95      17.35      0.20      -0.95
##      skew.2SE      kurtosis      kurt.2SE      normtest.W      normtest.p
##      -7.56      4.36      17.44      0.78      0.00
```

```
round(quantile(data$prate, probs = c(1, 5, 10, 25, 50, 75, 90, 95, 99,
                                     100)/100), 1)
```

```
##      1%      5%      10%      25%      50%      75%      90%      95%      99%      100%
##      31.8      54.0      62.8      78.1      95.7      100.0      100.0      100.0      100.0      200.0
```

```
data$prate[data$prate > 100]
```

```
## [1] 200.0 177.2 200.0
```

Based on the R output above, **prate** is one of the 8 variables contained in the dataset. There are **1534** observations of **prate**, and 0 of them correspond to NA values.

Its minimum and **maximum** values are 3.0 and **200.0**, respectively: the latter must correspond to an error, since a rate cannot be greater than 100%. A further analysis reveals that there are **3** observations in which **prate** exceeds 100, so we should and shall **discard** them from our analysis.

Its **mean** and **median** values are 87.6 and 95.7, respectively.

The last two values in the output of the `stat.desc()` function used above show the results of a Shapiro-Wilk test, which indicates non-normality (but that test is not very conclusive with such a large sample).

The **excess kurtosis** (the kurtosis minus 3) is positive (4.36), which indicates that the distribution of the sample is **leptokurtic** (a more acute peak around the mean and fatter tails than the normal distribution). Its skewness is negative (-0.95), which indicates that the distribution of the sample is **left-skewed** (it has a long left tail). The following Figures confirm both aspects.




Figure 1: Histogram of participation rate (%) in 401K plans of a company's employees (bin width = 5)

Though hard to see (roughly 44.3% of the observations correspond to the exact value of 100%), the histogram above reveals the one observation with a value of 177.2 and the two observations with a value of 200.0.




Figure 2: Approximate density plot of participation rate (%) in 401K plans of a company's employees

If we omit the 3 anomalous observations (which we'll do from now on), the (approximate) density plot looks like this (**not normal** at all: neither it is symmetrically distributed about the mean (87.4) nor it can take any positive or value):




Figure 3: Approximate density plot of participation rate (%) in 401K plans of a company's employees (excluding wrong values, higher than 100%)

Question 2

Your independent variable will be `mrate`, the rate at which a company matches employee 401K contributions. Examine this variable and comment on the shape of its distribution.

```
# First, discard anomalous observations of prate
data2 <- data[data$prate <= 100, ]
# Descriptive statistics of prate
summary(data2$mrate)
```

```
##      Min. 1st Qu.  Median    Mean 3rd Qu.    Max.
##    0.010  0.300   0.460   0.732  0.830   4.910
```

```
round(stat.desc(data2$mrate, desc = TRUE, basic = TRUE, norm = TRUE), 2)
```

```
##      nbr.val    nbr.null    nbr.na      min      max
##    1531.00      0.00      0.00     0.01     4.91
##      range      sum    median     mean  SE.mean
##     4.90    1120.71     0.46     0.73    0.02
## CI.mean.0.95      var    std.dev   coef.var  skewness
##     0.04      0.61     0.78     1.07    2.59
##      skew.2SE    kurtosis   kurt.2SE  normtest.W  normtest.p
##     20.71      7.59     30.35     0.70     0.00
```

```
round(quantile(data2$mrate, probs = c(1, 5, 10, 25, 50, 75, 90, 95, 99,
                                       100)/100), 1)
```

```
##    1%    5%   10%  25%  50%  75%  90%  95%  99% 100%
##    0.0   0.1   0.2   0.3   0.5   0.8   1.7   2.4   4.1   4.9
```

As with `prate`, there are no NA values of `mrate`.

Its minimum and maximum values are 0.0 and 4.9 (these values, as those of `prate`, correspond to percentages).

Its **mean** and **median** values are 0.732 and 0.460, respectively.

Again, a Shapiro-Wilk test (see the last two values in the output of the `stat.desc()` function) indicates that the distribution is far from normal, but this kind of test is not conclusive whatsoever with such a large sample.

The **excess kurtosis** (the kurtosis minus 3) is positive (7.59), which indicates that the distribution of the sample is **leptokurtic** (a more acute peak around the mean and fatter tails than the normal distribution). Its skewness is positive (2.59), which indicates that the distribution of the sample is **right-skewed** (it has a long right tail). The following Figures confirm both aspects.




Figure 4: Histogram of companies' match rate (%) to their employees' 401K contributions (bin width = 0.1)




Figure 5: Approximate density plot of companies' match rate (%) to their employees' 401K contributions

Question 3

Generate a scatterplot of `prate` against `mrate`. Then estimate the linear regression of `prate` on `mrate`. What slope coefficient did you get?




Figure 6: Scatterplot of the participation rate (%) to 401K plans of a company's employees against the match rate (%) of that company to their employees' contributions

While a low match rate of a company may correspond to almost the whole range of employees' participation rates, higher match rates correspond to high participation rates, which seems to indicate the positive relationship between both variables.

```
params <- "mrate" # regressor()
model <- lm(as.formula(paste("prate", paste(params, sep = "",
                                         collapse = " + "), sep = " ~ ")),
            data = data2)
summary(model)
```

```
##
## Call:
## lm(formula = as.formula(paste("prate", paste(params, sep = "",
##      collapse = " + "), sep = " ~ ")), data = data2)
##
## Residuals:
##      Min       1Q   Median       3Q      Max
## -82.289  -8.200   5.186  12.723  16.821
##
## Coefficients:
##              Estimate Std. Error t value Pr(>|t|)
## (Intercept)  83.0618     0.5641  147.24  <2e-16 ***
## mrate         5.8623     0.5275   11.11  <2e-16 ***
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
##
```



```
## Residual standard error: 16.09 on 1529 degrees of freedom
## Multiple R-squared:  0.07475,    Adjusted R-squared:  0.07414
## F-statistic: 123.5 on 1 and 1529 DF,  p-value: < 2.2e-16
```

Table 1: Effect of a company match rate to 401K plans on its employees' contribution

Employees' participation rate (%) to 401K plans	
Company match rate (%)	5.862*** (0.527)
Baseline (Intercept)	83.062*** (0.564)
R^2	0.075
F	123.525
p	1.2e-27
N	1531

As shown in [Table 1](#) above, **the slope coefficient is 5.862 (0.527)**: a 1 percentage point increase in the match rate would correspond to an increase of almost 6 percentage points in the participation rate to 401K plans of the employees, which could indicate that more employees are willing to make this kind of investment when their companies promote 401K plans by matching a higher percentage of their own contributions.

Question 4

Is the assumption of zero-conditional mean realistic? Explain your evidence. What are the implications for your OLS coefficients?

One of the ways to check this assumption is by plotting the residuals (\hat{u}) against the fitted values (\hat{y}), or even against the regressor when only one is used. That plot (below) shows that the residuals change with the fitted values (see the blue line; the smoother is far from flat), which suggests that the assumption may **not** be **realistic** (as well as a possible non-linear relationship between `mrate` and `prate`).




Figure 7: Scatterplot of fitted values of the regressand against the residuals

The implication is that the OLS coefficients are **biased**. But though the assumption of zero-conditional mean ($E(u|x) = 0$) is not met, the assumption of zero mean ($E(u) = 0$) and zero correlation ($Cov(x, u) = 0$) seems realistic (i.e., the **exogeneity** assumption is met), so they may be **consistent** at least:

```
mean(model$residuals)
```

```
## [1] -1.25493e-15
```

```
cov(model$residuals, data2$mrate)
```

```
## [1] 2.231346e-15
```

Question 5

Is the assumption of homoskedasticity realistic? Provide at least two pieces of evidence to support your conclusion. What are the implications for your OLS analysis?

The funnel shape of Figure 7 suggests that the assumption of homoskedasticity is **not realistic**, based on the shape of the scatterplot: the variance of the residuals is much greater for lower values of \hat{y} than it is for greater values.

The fact that about 80% of the data points in Figure 7 are concentrated in about one fifth of the x-axis (corresponding to `mrate` < 1) may make it look like there's more variation than is actually occurring. But there are more evidence of homoskedasticity...

A second piece of evidence comes from the analysis of the scale-location plot:




Figure 8: Scale-location plot of the regression model

The smoother (the blue line in the figure above) should be flat, which is not the case.

We might also run a Breusch-Pagan test, but with such a large dataset it is very likely that the result will be significant even if the assumption of homoskedasticity were realistic.

```
bptest(model) # Breusch-Pagan test

##
##      studentized Breusch-Pagan test
##
## data:  model
## BP = 27.921, df = 1, p-value = 1.264e-07
```

The implications are that:

1. we cannot conclude that the OLS estimators have the smallest variance among *all linear*, unbiased estimators, and
2. we cannot compute the variance of our estimates of the coefficients, so we don't have a clear idea of their precision.

Question 6

Is the assumption of normal errors realistic? Provide at least two pieces of evidence to support your conclusion. What are the implications for your OLS analysis?




Figure 9: Q-Q plot of the residuals of the regression

The left tail of the Q-Q plot twists counterclockwise, while the right tail twists clockwise, which indicates that the distribution is left-skewed (and not normal).

Let's also check some descriptive statistics of the residuals. As seen below, the **excess kurtosis** (the kurtosis minus 3) is positive (2.28), which indicates that the distribution of the sample is **leptokurtic** (a more acute peak around the mean and fatter tails than the normal distribution), and the skewness is negative (-1.44), which indicates that the distribution of the sample is **left-skewed** (it has a long left tail). Both details indicate a non-normal distribution.

```
format(stat.desc(model$residuals, desc = TRUE, basic = TRUE, norm = TRUE),
       digits = 3, drop0trailing = TRUE, scientific = TRUE, trim = TRUE)
```

##	nbr.val	nbr.null	nbr.na	min	max
##	"1.53e+03"	"0.00"	"0.00"	"-8.23e+01"	"1.68e+01"
##	range	sum	median	mean	SE.mean
##	"9.91e+01"	"-1.92e-12"	"5.19"	"-1.25e-15"	"4.11e-01"
##	CI.mean.0.95	var	std.dev	coef.var	skewness
##	"8.06e-01"	"2.59e+02"	"1.61e+01"	"-1.28e+16"	"-1.44"
##	skew.2SE	kurtosis	kurt.2SE	normtest.W	normtest.p
##	"-1.15e+01"	"2.28"	"9.14"	"8.57e-01"	"3.73e-35"

Let's plot the approximate distribution of the residuals:

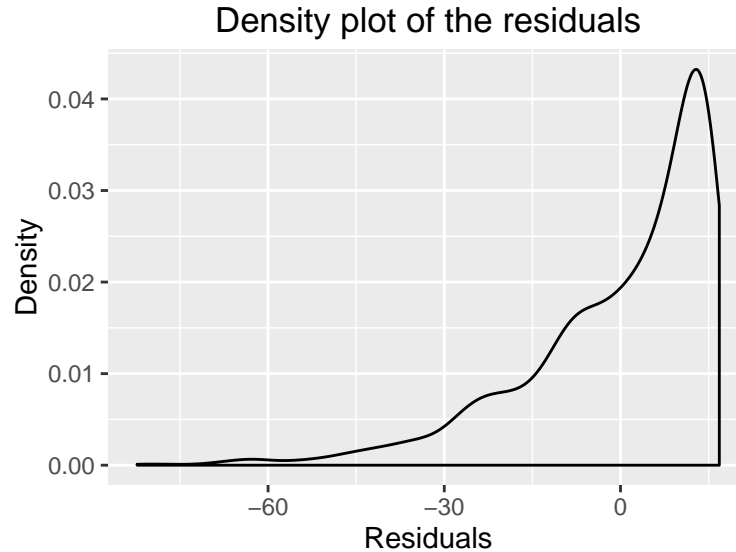


Figure 10: Density plot of the residuals of the regression

As expected, the density plot of the residuals has a shape similar to the explained variable; see [Figure 3](#).

Finally, the last two values in the output of the `stat.desc()` function used in the previous page showed the results of a Shapiro-Wilk test, which indicates non-normality (but that test is not very conclusive with such a large sample). We can confirm that result using the specific R function:

```
shapiro.test(model$residuals)
```

```
##
##  Shapiro-Wilk normality test
##
## data:  model$residuals
## W = 0.85699, p-value < 2.2e-16
```

The implication of non-normality of the errors is two-fold:

1. the OLS estimators will not have the smallest variance among *all* (linear or not) unbiased estimators (assuming the Gauss-Markov assumptions are met, which is not the case here), and
2. the t and F statistics do not have t and F distributions, respectively, which makes hypothesis testing impossible.

But if the Gauss-Markov assumptions were met (which is not the case; for example, the zero conditional mean and the homoskedasticity assumptions are broken), we could still use assume that the (standardized) OLS estimators are *asymptotically normally distributed*. I.e., we might use the t_{n-k-1} distribution for hypothesis testing, since our sample is quite large (1531 observations after discarding the 3 ones): the actual significance level we may get will be close to the one we chose (5% or any other).

Question 7

Based on the above considerations, what is the standard error of your slope coefficient?

The standard error of the slope coefficient (which is 5.862) is shown in [Table 1](#): **0.527**.

```
summary(model)$coefficients[2, 2]
```

```
## [1] 0.5274588
```

But the variance of the residuals is not constant, so we should use **robust standard errors** instead (which is actually smaller for the slope coefficient: **0.470**).

```
coeftest(model, vcovHC(model))[2, 2]
```

```
## [1] 0.470155
```

Table 2: Effect of a company match rate to 401K plans on its employees' contribution (using heteroskedasticity-robust SEs)

Employees' participation rate (%) to 401K plans	
Company match rate (%)	5.862*** (0.470)
Baseline (Intercept)	83.062*** (0.613)
R^2	0.075
F	155.471
p	4.7e-34
N	1531

Question 8

Is the effect you find statistically significant, and is it practically significant?

As shown in [Table 2](#), the effect of `mrate` on `prate` is highly **statistically significant**: the exact p -value is $p = 4.7e - 34$ (lower than what we got in [Question 3](#): according to [Table 1](#), $p = 1.2e - 27$ when we do not use heteroskedasticity-robust standard errors).

If a company matches one percentage point more of its employees' contributions, the contribution of those employees to 401k is, on average, about 5.9 percentage points higher, which is also a **practically significant** effect.